

On the Security and Privacy of Hyperledger Fabric: Challenges and Open Issues

Sotirios Brotsis*, Nicholas Kolokotronis*, Konstantinos Limniotis*, Gueltoum Bendiab†, and Stavros Shiaele†

*University of Peloponnese, 22131 Tripolis, Greece

Email: {brotsis, nkolok, klimn}@uop.gr

†University of Portsmouth, PO1 2UP, Portsmouth, UK

Email: {gueltoum.bendiab, stavros.shiaele}@port.ac.uk

Abstract—In the last few years, a countless number of permissioned blockchain solutions have been proposed, with each one to claim that it revolutionizes the way of the transaction processing along with the security and privacy preserving mechanisms that it provides. Hyperledger Fabric is one of the most popular permissioned blockchain architectures that has made a significant impact on the market. However, there are only few papers of finding architectural risks regarding the security and the privacy preserving mechanisms of Hyperledger Fabric. This paper separates the attack surface of the blockchain platform into four components, namely, consensus, chaincode, network and privacy preserving mechanisms, in all of which an attacker (from inside or outside the network) can exploit the platform's design and gain access to or misuse the network. In addition, we highlight the appropriate counter-measures that can be taken in each component to address the corresponding risks and provide a significantly secure and enhanced privacy preserving Fabric network. We hope that by bringing this paper into light, we can aid developers to avoid security flaws and implementations that can be exploited by attackers but also to motivate further research to harden the platform's security and the client's privacy.

Index Terms—Hyperledger Fabric, cyber-security, consensus protocols, chaincode risks, network threats, privacy

I. INTRODUCTION

Hyperledger Fabric [1] (for simplicity Fabric), has recently obtained massive popularity with hundreds of implementations around the world, since it is quite scalable, lenient against faults, and robust. For these reasons, among others, it can satisfy more than enough and sufficiently better than any other permissioned blockchain solution the purposes of an enterprise-based environment. Only in the past year, it is considered as the most deployable distributed ledger, in various areas of interest, such as the IoT ecosystem [2], the supply chain finance [3], the medical data management [4] and more.

By its permissioned nature, Fabric is a closed system, in which only the participants that have obtained the necessary credentials are able to read or write to the ledger. These participants are called peers and only a subset of them can

approve transactions and in order to do so, they have to mention their identity along with their signatures [5]. This setting makes it easier for the peers to manage the transactions on the ledger and it is typically the reason why Fabric is much faster than the ongoing permissioned blockchains.

For the maintenance of the identities of all the participating nodes (clients, *ordering service nodes* (OSNs) and peers) responsible is the membership service provider, which is one the most critical components of the platform, since it manages any type of access by issuing credentials in a form of cryptographic certificates that are used for authentication and authorization.

The features of Fabric are not limited only to its design, the support of pluggable consensus, which is another critical component; provides an unprecedented level of extensibility and specific the support of multiple ordering nodes that establish consensus regarding the transactions' total order. Moreover, since version 1.0, Fabric's ordering service comes without any *Byzantine Fault Tolerant* (BFT) consensus protocol that can address possible malicious ordering nodes, implementing only *Crash-Fault-Tolerant* (CFT) protocols based on Kafka and on Raft.

Nearly almost all the permissioned blockchain solutions can implement smart contracts, which are based on a programmable application logic that is being called each time a transaction is being proposed. In Fabric's case the smart contracts are realized by means of an arbitrary program that is authored in Go; the chaincode. The chaincode is executed by a set of peers locally and before each transaction is appended into the ledger, an output of the chaincode's execution is taken into account in order to decide whether a transactions is valid or not and which data will be included to the ledger.

From the security perspective, we analyze Fabric into four interconnected components, in which possible attacks and leakage of private information can occur; namely: the consensus, the chaincode, the network and its privacy preservation mechanisms. As a result of our research, in each of these components, we define the risks and we provide the corresponding counter measures needed to be addressed and enhance its security: a) The implemented consensus protocols in Fabric can withstand only some of the ordering nodes to crash but not to behave maliciously. b) The privacy protection



This project has received funding from the European Union's Horizon 2020 research and innovation programme under grant agreement no. 786698. The work reflects only the authors' view and the Agency is not responsible for any use that may be made of the information it contains.

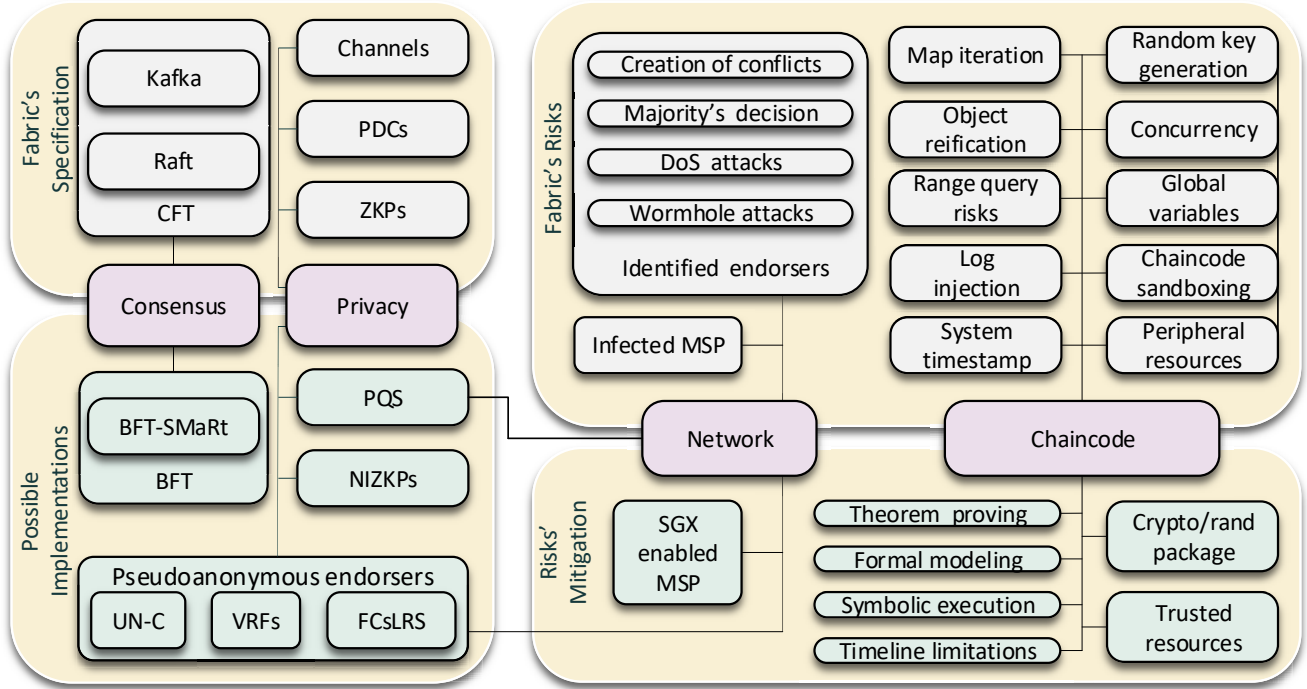


Fig. 1. Hyperledger Fabric's security architecture

features of Fabric include the following four aspects, namely the channels, the private data collections, the zero-knowledge proofs and the membership service provider, but the platform does not support any non-interactive zero-knowledge proofs and post quantum signatures. *c)* and *d)* Regarding the security of the smart contracts (and the network's) we have surveyed and outlined the potential vulnerabilities associated with Fabric's chaincode (and the network's layer) and provided the corresponding countermeasures and verification methods for these vulnerabilities to be averted, as shown in Figure 1. To the best of our knowledge there is no other published research that has assembled from all the four components, the possible threats that can occur in Fabric and has discussed possible counter-measures to mitigate these threats.

The remainder of this paper is organized as follows. Section II presents the security that Fabric provides from the consensus process and a BFT consensus protocol that can be implemented in the presence of malicious OSNs along with a comparison of these protocols. Section III surveys a number of smart contracts' vulnerabilities and the possible counter measures and methods that can be taken to enhance and verify the chaincode's security. In Section IV we analyze the impact to the network; if the membership service provider is compromised, the possible attacks that can occur from having the endorsing peers been identified and we propose possible countermeasures. In Section V, we analyze how Fabric satisfies many privacy flavors that and we discuss the open challenges to these issues.

II. CONSENSUS SECURITY

Consensus in blockchain technology is a state, in which all the participants agree upon an identical sequence of the messages in a specific order [1], i.e. the atomic broadcast [6]. In Fabric the OSNs atomically broadcast endorsements to establish agreement, by means of consensus, regarding the transactions' total order.

A. Consensus methods and Fabric's specification

The support of pluggable consensus protocols is a valuable discrimination in each blockchain protocol to achieve a trust model or other use cases [1]. Assuming that a blockchain protocol is deployed in a small environment, maintained by a small enterprise or a single trusted-third party; the implementation of complicated protocols can be seen as an exaggerated drag on the energy's consumption. For this reason, the adoption of a single OSN to achieve consensus and order the clients' proposals in a small and highly trusted environment can be proved to be a quite as good choice as the adoption of other perplexing consensus protocols.

1) *Solo*: As a centralized consensus protocol, Solo is being implemented for testing purposes, since it necessitates only a single node to receive and order the incoming transactions in order to provide an inferior consensus method. By adopting this approach, developers can focus on other concerns such as the development and improvement of the chaincode. Nevertheless, Solo's outcome is the creation of a *Single-Point-of-Failure* (SPoF) and if this single ordering node crashes the entire system will collapse.

2) *Apache Kafka*: Kafka [7] is a distributed publish/subscribe messaging pattern that is being used to transfer large amount of log data with significantly low latency. Kafka's key components are the producers, the topics, the consumers and the brokers. The recorded information is published from the producers to a stream of messages called the "topic", which is a partition of segments of files. The messages are stored from the brokers as the latest segment file and when the producers publish messages to the partitioned logs, only the subscribed consumers can consume these messages - sequentially - by making requests to the brokers. Kafka's fault tolerance properties derive from the ZooKeeper [8] and they can be achieved with the replication of the partitions between the brokers. Since, Kafka follows the leader-follower design, each partition has its own leader, whose actions are passively replicated by its followers. Therefore, if the ongoing leader crashes, a new election process begins with one of its followers to take its place. Regarding its performance, Kafka showed momentous results in [1] and brought significant enhancements in the area of the business oriented and permissioned blockchains. Thus, Kafka, surpassing Solo, is the suggested protocol in version 1.0.

3) *Raft*: The Raft consensus protocol [9] is also based on the "leader-election" model to establish consensus by electing a leading node to acquire the incoming entries from the clients and replicate them. To provide strong leader and coherency, the protocol is separated into three phases, i.e. the leader election, the log replication and the safety. The time in Raft proceeds in arbitrary time periods, called "terms", with each term to be defined by an increasing number. The nodes in Raft are hierarchically ranked in different states, with each node to be either a leader, a follower or a candidate. The leader is the principal entity of the protocol and it is elected per channel, with the task to interact with the clients and then replicate its entries to its synchronized followers. Therefore, in order to achieve the best possible synchronization, it sends systematic heartbeats to its followers and even if the network suspects that the leader has crashed, at least one of its followers will detect this divergence, cast a vote to the network and attempt to take its place [10]. Some nodes might compete to win the election by seeking votes from other nodes. Therefore, these nodes are considered as "candidates".

Raft ensures that exclusively a single node can become a leader, even if some nodes miss a term or split ups occur. In the first case the outdated node will revise the term's number and fall into the follower's state; and in the latter, the current term will be ended without any outcome of the election process. Raft's performance is not fully tested, but its adoption as the recommended consensus protocol since the version 1.4.1, showed that Raft can provide thousands of transactions in real world scenarios [11] and twice the transaction's throughput of Kafka with even less latency [12].

4) *BFT-SMaRt*: BFT-SMaRt [13] is java-based consensus protocol that can provide a secure ordering service for Fabric [14]. In the ideal case where no adversarial *validation replicas* (VRs) exist in the network, the BFT-SMaRt's message

processing is identical to PBFT's [15].

WHEAT is the component that BFT-SMaRt's ordering service [14] relies upon [16] to provide a powerful vote assignment scheme, low latency and fast replication among the VRs, without imperiling the network's security. The ordering service consists of the cluster nodes and the frontends. By adopting BFT-SMaRt as the consensus protocol, the transaction flow is almost the same to [1]. Upon the collections of the endorsements from the peers, a client generates a signed envelope that contains the channel's identity and the peer's endorsements accompanied by their signatures [1]. This envelope is disseminated to the frontends and finally sent to the OSNs for ordering. When the OSNs collect a predetermined number of envelopes from some trusted frontends or when a predetermined time has elapsed, a new block is created that contains only the valid transactions. Therefore, the transactions' validation differs from [1] and it occurs prior to the creation of the block's signature. Finally, all the same, the block is transferred to the frontends and then to the peers that manage the ledger. BFT-SMaRt's performance is measured in [14] and showed that it can achieved throughput more than 10 thousand TPS and the block confirmation time to be less than 1sec.

B. Comparison

The adoption of Solo creates a SPoF, since it is not lenient against malicious failures and crashes and for this reason, it should not be deployed in real world environments. Both Raft and Kafka benefit from the "leader-follower" model to address crashes, but despite Kafka's popularity, various intricate components have to be handled for its implementation. In Raft's case, these components are enclosed to the ordering service [11], meaning that there are less components that might crash. Kafka is designed to be deployed in an environment with a small number of OSNs and the cluster to be managed by a single entity. This concept does not contribute much to escalate the system's decentralization, forcing all the OSNs to be ulcerated by a single entity. In contrast to Kafka, a Raft-based ordering service is more decentralized, more scalable and it can achieve a greater throughput [12]. For reasons as such, the Kafka and the Solo consensus protocols are deprecated in Fabric's version 2.0. Despite Kafka's deprecation though, a reconstruction of the Fabric's transaction flow managed to increase Kafka's throughput from 3.5 to 20 thousand TPS [17].

On the other hand, BFT-SMaRt can not only achieve a remarkable transaction performance of 10 thousand TPS but it also can withstand tolerance against possible malicious behavior by the OSNs. In BFT-SMaRt, the invalid transactions are not included in the blocks, since the transactions' validation occurs prior to the block's creation and dissemination to the peers. Although BFT-SMaRt seems to outperform Raft, both in terms of security and performance; it is a Java-based library that it does not provide a very stable ordering service and therefore it is not being currently adopted by the Hyperledger Project. BFT-SMaRt's drawback though is that it necessitates two processes to operate; the first built in Java and the second

from GO, using a network socket bound among them, which can act as bottleneck in the system's operation.

C. Consensus' challenges and open issues

The consensus protocol is the most critical component of a distributed ledger. CFT protocols are considered to be contradicting to the platform's security, since any malicious action can affect the network's security. The BFT protocols are adopted to permissioned blockchain solutions for targeted use cases, in which the requirements to provide a secure implementation is more evident. The fact that each peer in a permissioned network is accountable for its behavior, provides an incentive to the nodes to adhere to the protocol. In spite of the tolerance that the consensus protocol provides (i.e. CFT or BFT), Fabric can mitigate the most familiar and sophisticated consensus-oriented attacks, such as the double spending, by its design. Therefore, the consolidation of Fabric and a CFT consensus protocol can be considered as an ideal implementation in a confidential network, such as an enterprise environment.

III. SMART CONTRACTS' SECURITY

The execution of the smart contracts defines operations that have been acknowledged by all the participating entities; with each one to execute the contract locally, propose the result to the network and then collaborate with it to select which result is going to be inserted into the ledger. The development of smart contracts can incorporate several programming errors that can ultimately lead to exploitable bugs and malicious behavior. Consequently, the smart contracts are prone to code errors and inconspicuous vulnerabilities, while their accuracy and security can be violated by malicious programmers by means of exploits.

A. Chaincode' vulnerabilities and Fabric's specification

The programs' defects, such as coding flaws and designing errors, are the main reasons that cause the smart contracts' vulnerabilities. In [18], a large variety of Ethereum smart contracts vulnerabilities have been identified and showed that programming flaws lead to faulty behavior, a series of attacks and possible exploits. As we are focused on Hyperledger Fabric and particularly to the Fabric's chaincode, our following discussion is focused on risks that might ascend from the programming languages, the Fabric's features and the misunderstanding of the common practices.

Fabric's smart contracts were not built with the vision of being strained to a *domain-specific programming language* (DSL), such as the Ethereum's Solidity, but rather to be authored in high-level languages (such as Go, Java and Node.js) to reduce the developers' learning cost. Formally, the smart contracts that have been authored in DSLs are ruled by particular features and restrains for blockchains. Since this is not the case of Fabric, the known risks and vulnerabilities might differ with the risks associated with general-purpose programming languages. Therefore, based on [19], [20], [21] and [22], we outline in the following table the

most prominent security vulnerabilities in Fabric's chaincode. The programming language Go is the most widely used during the chaincode's development and thus, most of these security vulnerabilities derive from its non-deterministic behavior. Nevertheless, Fabric has no native cryptocurrency built in and thus, it is not easy to define how severe each vulnerability is and how easy or difficult it is for an attacker to exploit possible bugs and execute double spending attacks.

In the execution phase, the peers do not execute the chaincode at the same time and in the same environment. If a transaction is valid and the results that derive from the chaincode's execution are not deterministic, in the validation phase these results might be rejected or allow double spent transactions to be included to the ledger. The risks that derive from the Go language, from the platform's features or from mis-understanding of the common practices can lead to inconsistencies to the peers' ledger:

1) *Random key generation*: Since, the simulation of the chaincode occurs in each endorsing peer, the random seed for the keys is different. In Go, the random seed is set as 1 and therefore can be easily predicted by a client.

2) *Object reification*: The value of the variables are handled through a pointer, which is an address of memory. Therefore, using reified object addresses might cause non-determinism.

3) *System Timestamp*: It is difficult to ensure that the timestamps are run concurrently in each peer.

4) *The global state variables*: Global state variables that are not stored to the ledger, might change innately and cause inconsistencies to the endorsing results.

5) *Concurrency*: If multiple transactions are executed concurrently and under high load, a possible change at the keys' versions might lead to key collisions and double spending. For example, if a transaction has passed from the endorsing phase and its version key has changed before it reaches the validation phase not only a program error will occur but this action might also allow another (possible double spent) transaction to be included to the ledger.

6) *The non-determinism that ascends from peripheral resources*: Some accessing resources, such as web services, external libraries, external files and system command executions; can corrupt the code and return different endorsing results among the endorsing peers.

7) *Range query risks*: Queries methods to access the Fabric's state databases and obtain private data, (e.g. the history or the state of a key); are not executed again in the validation phase and can lead to phantom reads, in which the dirty data cannot be detected.

8) *Chaincode sandboxing*: Although, Fabric's chaincode is executed in an isolated docker container and provides sufficient privileges, it can be exploited in a malicious way, such as to execute port scanning, identify and exploit vulnerable peers, install malicious software and execute attacks.

9) *Log injection*: Any corruption of the log messages can possibly avert them from being executed automatically and allow the attacker to view the processed logs.

10) *Map structure iteration*: Due to the hidden implementation details of the Go programming language, when an iteration with map structure is used, non-determinism may arise and the order of key values might be different.

B. Chaincode's challenges and open issues

A docker container, by its design, is significantly secure, particularly when the processes are executed with no privileged users in the container. Albeit, developers should be familiar with the dockers' security issues to fully take advantage of the security and the efficiency that the dockers provide. Thus, in this section we discuss the proactive measures that can be reserved to increase the smart contracts' security:

- Timeline limitations of the chaincode's execution can be employed;
- The chaincode's execution can be performed with privileges other than root;
- GO's crypto/rand package can be used to produce a cryptographically secure random seed for the keys and minimize the risk of double spending;
- Any changes at the keys' versions can increase the possibility of double spending and thus, should be avoided;
- Only trusted services peripheral of the platform should be accessed;
- Chaincode's input arguments should be checked for escape characters;

The aforementioned counter-measures can be easily taken into account during the chaincode's implementation. Nonetheless, a number of other techniques [20] are relied upon formal verification approaches to mitigate risks in smart contracts:

1) *Theorem proving*: One of the most ordinary methods to formalize smart contracts. Symbolic logic (with axioms or premises) is used in theorem proving to prove the necessary security properties that enhance the contracts' accuracy.

2) *Symbolic execution tools*: This technique uses testing tools, such as the Chaincode Scanner, to perform static analysis and find bugs, vulnerabilities and bad practices in the chaincode, as well as, to provide a detailed description of a possible problem.

3) *Formal modeling*: This method formalizes the smart contracts' risks, by using precise statements that define the relationship between the smart contracts' components resulting in a) unambiguous communication, and b) replicable/reproducible results. These risks are classified into seven categories [23], that analyze the platform's properties; specifically: "privacy, security, bug bounty, trustworthiness of data, scalability and correctness of the consensus protocol.

Although, these techniques are not the most common for verification of the smart contracts (according to [23]), they seems to be a quite promising field of research, especially in the Fabric's case [24].

IV. NETWORK'S SECURITY

With the increase of interest in permissioned blockchains; networks that possess the qualifications to handle administrative permissions, the possibility of an anonymous attacker

feigning to behave correctly is almost minimized. Attacks and strategies such as the 51%, the Sybil attack, the block withholding and the selfish strategy are lesser risks to the network's security [25], [26], due to the trust and the restricted permissions to those given access to the network.

A. Network's threats and Fabric's specification

The rise of Fabric, on the other hand, comes with new security risks and concerns that can harm the network's operation and performance [27]. Therefore, in this section, we present the most prominent attacks that derive from the network's level, if some participating entities are compromised; along with feasible countermeasures to enhance its security.

1) *Compromised membership service provider (MSP)*: The most revolutionary aspect of blockchains lies in the decentralized nature of a trustless network and Fabric seems to violate just that. The centralized aspect of Fabric lies on the MSP and the following *Certificate Authority (CA)*. The MSP is a critical component of the platform, since it manages the registration, the identities and the type of access of all the nodes in the network; comprising the clients, the peers and the OSNs. Therefore, if the MSP is compromised, administrative controls such as adding and removing identities to and from the network, as well as, the type and the amount of the given access to the existing nodes are all managed entirely by the attacker. With a malicious MSP, the unauthorized access that is given to the attacker can cause serious damage and possible lead to further attacks such as: invalid identification attack, double-spending, attacks on the CA, etc. [27].

2) *Identified endorers*: In Fabric's execution phase, a transaction needs to be approved by a set of endorers. In return, the endorers, approve the transaction mentioning their identity along with their signatures so that they can be verified in the forthcoming phases. However, Fabric has some drawbacks when the peers are identified [28], [29], namely:

- Creation of conflicts: For some transaction, the endorsing peers may have different opinion about their validity. By revealing their identities, possible conflicts can be created within the consortium [29].
- Majority's decision: The endorsement policy not only does not allow the endorers to approve the transaction in secret, but also takes into consideration the majority's decision.
- DoS attacks: An opening for DoS attacks on selective endorers might be created, either to halt specific transactions from being included to the ledger or to degrade the network's performance [28].
- Wormhole attacks: An opening for wormhole attacks might be created, if a peer behaves maliciously and colludes with an adversary exterior to the channel. This attack can lead to leakage of private information of all the peers of a specific channel. This problem is concluded by defining the fact that the Fabric's access control mechanism is depended on trusting each peer inside a channel [30], [31].

B. Network's challenges and open issues

The aforementioned analysis for the MSP and the endorers indicates that the appropriate safeguards are not in place to mitigate network's risks in the context of Fabric. Although the platform offers strong accountability in the network to meet the necessary flavors of security, the research in this field is still ongoing. Thus, some other possible approaches that could be appropriately adopted in the (near) future are:

1) *Securing the MSP*: The possible threats that can derive from a compromised MSP are tackled in a recent study [32] with the adoption of *Intel Software Guard Extensions* (SGX). The SGX remote attestation techniques and the contained execution features that this method provides, can register each entity of the system as a trusted node. The capabilities that the SGX provides can secure the MSP in all the phases that it is invoked; including each node's registration, the transaction's signature and its verification. An SGX enabled MSP can also mitigate various privacy risks in Fabric and enhance the network's defense against possible attacks. Therefore, the attack surface of the membership service is required to be analyzed completely and rigorous proofs should be created to formally quantify all the risks that are aligned with it.

2) *Pseudonymizing the endorers*: Motivated by the aforementioned risks regarding the endorers identification, the authors in [29], created a ring signature scheme, named *Fabric's Constant-Sized Linkable Ring Signature* (FCsLRS) to pseudonymize the endorers' identities. They implemented this signature scheme in GO and provided experimental analysis about its security and performance by shifting the *Rivest-Shamir-Adleman* (RSA) modulus size. In the anonymized endorsement system that they have implemented, a threshold endorsement policy needs a set of endorers to approve a transaction without revealing their identity and only by counting and checking individual valid ring signatures. A similar work of [28], evaluated the outcome of a DoS attack on the endorers and proposed two anonymization techniques; the first by using *verifiable random functions* (VRFs) and the second by using pseudonyms. In both cases, there is a trade-off between the network's efficiency and its security; with both the sender and the receiver of a proposed transaction in a private channel to be anonymized [28]. However, this approach certified with security proofs of *Signature Unforgeability and Unlinkability in Ciphertext* (UN-C) that Fabric can be immune to DoS and wormhole attacks.

V. PRIVACY PRESERVATION

With the current emerge of the blockchain technology, the transactional data are being made credible (via consensus) and shareable (via the distributed ledger). However, the adoption of this advancement endangers the disclosure of the users' or the companies' private data. Actually, no user would want its private information to be revealed to unauthorized entities exterior to the network and no company would want its competitors to know any private information regarding the costs, the prices or the annual salaries of its employees. Therefore, the collection of the sensitive information and its

secure storage to the distributed ledger is a crucial issue, since it has to be complied with the *General Data Protection Regulation* (GDPR). Therefore, in this section, we present Fabric's privacy preservation mechanisms and the few but critical advancements that are needed to be made in order to further enhance the clients (users' or companies') privacy.

A. Privacy techniques and Fabric's specification

Fabric supports significant privacy protection mechanisms. Starting from its permissioned nature that authorizes the participating entities of the network to strongly authenticate their identity several features are provided to accommodate the necessary flavors of privacy.

1) *Channels*: A channel is a state partition with its own access policy rules and transactions' ordering mechanism. Each channel is managed by a set of peers and it is associated with some policies that provide access to the corresponding resources (such as, the ledger's state, the transactions included in it and the corresponding chaincodes). When a peer registers to a channel, which is characterized by a unique identifier; then the corresponding ledger is created and run on this peer allowing it to manage an identical and consistent data store with the rest of the channel's peers. Therefore, privacy-preservation mechanisms such as the channels, are highly important in cases of providing blockchain solutions into a consortium environment (where the consortium is comprised by a number of organizations or parties with common business goals).

2) *Private data collection*: While the channels are devoted to the preservation of the information's privacy, by allowing the information to be stored separately; the *Private Data Collection* (PDC) can preserve the privacy of data from another perspective [3]. The PDC allows, only a stated set of peers in a channel to preserve the actual data, while the remaining peers access only its existence proof. The PDC is created to provide to the peers the capabilities of endorsing, committing, or queering private data without being forced to create a new channel and add additional overhead.

The PDC is actually an accumulation of the following elements: 1) the actual private data; which are sent from/to authorized peers by means of a gossip protocol, stored on the peer's private state databases and can be seen only by this set of nodes and not by the OSNs (which are not involved in this case), and 2) the hash of private data that is executed, ordered, and stored on each peers ledger as evidence of the existence of the transaction. In some use cases, where a peer of a PDC wishes to share private data with other peers – for example, to transfer any asset to a *trusted third party* (TTP), the TTP can produce the hash of the private data and subsequently examine if the output of the hash value is consistent with the hash that is stored on the channel's ledger and thus, prove the existence of the transaction.

The "right-to-be-forgotten" can be used in the private transactions, since each peer can erase its own private database at any time, with the data itself to be deleted irreversibly and the hash pointing to the underlying private information to still

exist. An additional aspect that is implemented with the private transactions is the limitation of usage. A “BlockToLive” policy can be defined for each PDC to determine an amount of time that has to be elapsed with the concealed private database to be automatically erased.

Despite the privacy that PDCs provide, they should be used with caution, since the metadata of the private information, is much more than metadata and can be used to unlock the real private data. In this attack scenario, the unauthorized peers of the same channel can observe the shared ledger and detect if the private transactions occur periodically.

Concluding in this section, the obtained outcomes are that some peers will have full access to the ledger and others may only see what they are allowed to. In the case, where the transactional data must remain hidden during ordering from some peers of the same channel and the OSNs, the implementation of PDCs is the solution.

3) *Zero-Knowledge Proofs (ZKPs)*: ZKPs establish a significant cryptographic primitive to preserve and improve privacy in the blockchain platforms. In Fabric’s case, there are two privacy preserving mechanisms that are achieved with the implementation of ZKPs, and are:

- Identity mixer [33] that leverages ZKPs to provide to the clients anonymous authentication regarding their transactions’ proposals. The implementation of ZKPs might have a significant impact when a client’s actual identity and the attributes that it is associated with, must be kept secret from the network (such as, the peers). For example, if the peers wish to verify that a transaction is indeed sent from a correct client, which is either a member of a specific organization; (referred in [33] as “membership proof”), or it indeed possesses a specific set of attributes (also referred as “selective disclosure of attributes”). In both cases the identity mixer verifies that the client’s identity is not disclosed.
- *Zero-Knowledge Asset Transfer (ZKAT)* [33] is a method built on top of anonymous authentication mechanisms provided by the identity mixer that also uses ZKPs in various applications aiming at asset management along with audit support (referred as *Zero-Knowledge Asset Transfer (ZKAT)*). With this privacy-preserving mechanism, the clients can issue transactions without disclosing any other information to the peers regarding the exchange of assets; but only the evidence that each transfer is complied with the asset management rules.

B. Privacy-preservation challenges and open issues

The aforementioned analysis indicates that appropriate safeguards are in place to mitigate privacy risks in the context of Fabric. Although Fabric offers strong privacy-preserving mechanisms to meet the necessary flavors of privacy, the research in this field is still ongoing. Thus, some other possible approaches that could be appropriately adopted in the (near) future are:

1) *Non-interactive ZKPs*: Although Fabric prevents unauthorized peers to access channel resources, the transactional

data is disclosed to all the channel peers. This limitation can be overcome with FabZK [34]. FabZK is a proposed extension for Fabric to support auditable privacy-preserving chaincode by means of verifiable and well-structured cryptographic primitives, pertaining *non-interactive Zero-Knowledge Proofs* (NIZKPs) on Pedersen commitments. The proposed protocol, manages a set of APIs for the client code and the chaincode to establish automated validation, while it improves the transactions’ performance with two validation steps, in which each party executes active and lightweight auto-validation. FabZK showed significant results in [34], with its cryptographic primitives to outperform other approaches (such as the zkLedger [35]) when the NIZKPs are generated and verified.

2) *Post-quantum signatures*: To ensure secure communication, Fabric relies on *Public Key Infrastructure* (PKI) for the digital signatures and the digital identities that are perilous to the operational security of its network. Moreover, the GDPR demands “consistent” methods to be employed and protect each user’s personal identifiable information. However, Fabric’s ecosystem is not post-quantum secure, making all the information that is disseminated over the network to be vulnerable to malicious decryption techniques by a large scale of quantum computers. Therefore, it is left to see, if post-quantum digital signatures are going to be implemented in the (near) future [36].

VI. CONCLUSIONS

Appropriate implementations and counter-measures to mitigate possible risks in Hyperledger Fabric are discussed in this paper.

From the consensus protocols security perspective, BFT consensus protocols, such as BFT-SMaRt are currently being researched, since they can provide a significant transaction throughput and tolerance against malicious OSNs. Moreover, BFT protocols have not yet been deployed in production environments; it remains to be seen whether and how such protocols are going to be adopted by the Hyperledger Project.

Regarding the smart contracts, a plethora of techniques have been discussed to harden the chaincode’s security. However, a quite-promising trend relies upon formal verification approaches like theorem proving, formal modeling and symbolic execution to mitigate the chaincode’s risks. With these approaches, some other aspects of Fabric, such as privacy, performance, and scalability can be analyzed thoroughly.

From the network’s security perspective, representative attacks were discussed against the network when the MSP is compromised and proactive solutions having been proposed as possible mitigation methods. An alternative/complementary option would be to employ a TEE, such as Intel’s SGX, to address insider threats and also DDoS attacks resulting from the manipulation or aversion of the chaincode’s execution. Techniques to mitigate wormhole attacks vary from those relying on the anonymization of the senders and recipients in the transactions inside a channel, to those employing group signature approaches. Due to the need for accountability in

Fabric, such solutions need to be further assessed possibly along with privacy preserving solutions.

Appropriate safeguards to implement the basic privacy requirements, are discussed to mitigate the privacy risks in the context of Fabric. The implementation of ZKPs can achieve anonymous client authentication with identity mixer, privacy-preserving exchange of assets with ZKAT, and the “right to be forgotten” can be efficiently implemented with PDCs. As the research in this area is rapidly evolving, other promising approaches have been identified that could be adopted in the (near) future, such as the implementation of NIZKP and post-quantum digital signatures. These two areas seem to be a quite promising for research regarding the privacy preservation mechanisms of Fabric.

REFERENCES

- [1] E. Androulaki *et al.*, “Hyperledger fabric: A distributed operating system for permissioned blockchains,” in *Proc. of 13th EuroSys Conf.*, ser. EuroSys '18, 2018, pp. 30:1–30:15.
- [2] S. Brotsis *et al.*, “Blockchain solutions for forensic evidence preservation in iot environments,” in *2019 IEEE Conference on Network Software (NetSoft)*, 2019, pp. 110–114.
- [3] C. Ma *et al.*, “The privacy protection mechanism of hyperledger fabric and its application in supply chain finance,” *Cybersecurity*, vol. 2, no. 1, pp. 1–9, 2019.
- [4] H. Tian, J. He, and Y. Ding, “Medical data management on blockchain with privacy,” *Journal of medical systems*, vol. 43, no. 2, p. 26, 2019.
- [5] N. Kolokotronis *et al.*, “On blockchain architectures for trust-based collaborative intrusion detection,” in *2019 IEEE World Congress on Services (SERVICES)*, vol. 2642-939X, 2019, pp. 21–28.
- [6] F. Cristian *et al.*, “Atomic broadcast: From simple message diffusion to byzantine agreement,” *Information and Computation*, vol. 118, no. 1, pp. 158 – 179, 1995.
- [7] J. Kreps, N. Narkhede, and J. Rao, “Kafka: A distributed messaging system for log processing.”
- [8] P. Hunt *et al.*, “Zookeeper: Wait-free coordination for internet-scale systems,” in *USENIX Annual Technical Conference*, 2010.
- [9] D. Ongaro and J. Ousterhout, “In search of an understandable consensus algorithm,” in *2014 USENIX Annual Technical Conference (USENIX ATC 14)*, 2014, pp. 305–319.
- [10] E. Bellini, P. Ceravolo, and E. Damiani, “Blockchain-based e-vote-as-a-service,” in *2019 IEEE 12th International Conference on Cloud Computing (CLOUD)*, 2019, pp. 484–486.
- [11] H. Fabric, “Hyperledger fabric’s documentation,” https://hyperledger-fabric.readthedocs.io/en/release-2.0/orderer/ordering_service.html.
- [12] Christopher Ferris, “Does Hyperledger Fabric perform at scale?” <https://www.ibm.com/blogs/blockchain/2019/04/doeshyperledgerfabric-performatscale/>.
- [13] A. Bessani, J. Sousa, and E. E. P. Alchieri, “State machine replication for the masses with bft-smart,” in *2014 44th Annual IEEE/IFIP International Conference on Dependable Systems and Networks*, 2014, pp. 355–362.
- [14] J. Sousa, A. Bessani, and M. Vukolic, “A byzantine fault-tolerant ordering service for the hyperledger fabric blockchain platform,” in *2018 48th Annual IEEE/IFIP International Conference on Dependable Systems and Networks (DSN)*, 2018, pp. 51–58.
- [15] M. Castro, B. Liskov *et al.*, “Practical byzantine fault tolerance,” in *OSDI*, vol. 99, no. 1999, 1999, pp. 173–186.
- [16] J. Sousa and A. Bessani, “Separating the wheat from the chaff: An empirical design for geo-replicated state machines,” in *2015 IEEE 34th Symposium on Reliable Distributed Systems (SRDS)*, 2015, pp. 146–155.
- [17] C. Gorenflo *et al.*, “Fastfabric: Scaling hyperledger fabric to 20,000 transactions per second,” in *2019 IEEE International Conference on Blockchain and Cryptocurrency (ICBC)*, 2019, pp. 455–463.
- [18] N. Atzei, M. Bartoletti, and T. Cimoli, “A survey of attacks on ethereum smart contracts (sok),” in *Principles of Security and Trust*, M. Maffei and M. Ryan, Eds., 2017.
- [19] K. Yamashita *et al.*, “Potential risks of hyperledger fabric smart contracts,” in *2019 IEEE International Workshop on Blockchain Oriented Software Engineering (IWBOSE)*, 2019, pp. 1–10.
- [20] H. Hasanova *et al.*, “A survey on blockchain cybersecurity vulnerabilities and possible countermeasures,” *International Journal of Network Management*, vol. 29, no. 2, p. e2060, 2019.
- [21] T. Kaiser, “Automated security analysis of chaincode,” https://static.sched.com/hosted_files/hgf18/18/GlobalForumtobias-kaiser.pdf.
- [22] Y. Huang *et al.*, “Smart contract security: A software lifecycle perspective,” *IEEE Access*, vol. 7, pp. 150 184–150 202, 2019.
- [23] A. Singh *et al.*, “Blockchain smart contracts formalization: Approaches and challenges to address vulnerabilities,” *Computers & Security*, vol. 88, p. 101654, 2020.
- [24] C. Agostino Ardagna *et al.*, “Trustworthy iot: An evidence collection approach based on smart contracts,” in *2019 IEEE International Conference on Services Computing (SCC)*, 2019, pp. 46–50.
- [25] N. Kolokotronis *et al.*, “Secured by blockchain: Safeguarding internet of things devices,” *IEEE Consumer Electronics Magazine*, vol. 8, no. 3, pp. 28–34, 2019.
- [26] B. Shala *et al.*, “Blockchain and trust for secure, end-user-based and decentralized iot service provision,” *IEEE Access*, vol. 8, pp. 119 961–119 979, 2020.
- [27] A. Davenport, S. Shetty, and X. Liang, “Attack surface analysis of permissioned blockchain platforms for smart cities,” in *2018 IEEE International Smart Cities Conference (ISC2)*, 2018, pp. 1–6.
- [28] N. Andola *et al.*, “Vulnerabilities on hyperledger fabric,” *Pervasive and Mobile Computing*, vol. 59, p. 101050, 2019.
- [29] S. Mazumdar and S. Ruj, “Design of anonymous endorsement system in hyperledger fabric,” *IEEE Transactions on Emerging Topics in Computing*, pp. 1–1, 2019.
- [30] E. Bellini, “A blockchain based trusted persistent identifier system for big data in science,” *Foundations of Computing and Decision Sciences*, vol. 44, no. 4, pp. 351 – 377, 2019.
- [31] E. Bellini, Y. Iraqi, and E. Damiani, “Blockchain-based distributed trust and reputation management systems: A survey,” *IEEE Access*, vol. 8, pp. 21 127–21 151, 2020.
- [32] X. Liang *et al.*, “Towards a trusted and privacy preserving membership service in distributed ledger using intel software guard extensions,” in *Information and Communications Security*. Springer International Publishing, 2018, pp. 304–310.
- [33] E. Androulaki, “Private and confidential transactions with hyperledger fabric,” <https://developer.ibm.com/technologies/blockchain/tutorials/cl-blockchain-private-confidential-transactions-hyperledger-fabric-zero-knowledge-proof/>.
- [34] H. Kang *et al.*, “Fabzk: Supporting privacy-preserving, auditable smart contracts in hyperledger fabric,” in *2019 49th Annual IEEE/IFIP International Conference on Dependable Systems and Networks (DSN)*, 2019, pp. 543–555.
- [35] N. Narula, W. Vasquez, and M. Virza, “zkledger: Privacy-preserving auditing for distributed ledgers,” in *15th USENIX Symposium on Networked Systems Design and Implementation (NSDI 18)*. USENIX Association, 2018, pp. 65–80.
- [36] R. Campbell, “Transitioning to a hyperledger fabric quantum-resistant classical hybrid public key infrastructure,” *The Journal of The British Blockchain Association*, p. 9902, 2019.